

 [alecjacobson](#) / [geometry-processing-mesh-reconstruction](#)

Mesh Reconstruction assignment for Geometry Processing course

 github.com/alecjacobson/geometry-processing

 MPL-2.0 License

☆ 43 stars 🍴 67 forks

☆ Star

👁 Watch ▼

<> Code

! Issues 5

🔗 Pull requests 34

🎮 Actions

📁 Projects

📖 Wiki

! Settings

🔗 master ▼

...



alecjacobson another bump? this time from local ...

22 hours ago

🕒 48

[View code](#)

README.md

Geometry Processing – Mesh Reconstruction

To get started: Clone this repository then issue

```
git clone --recursive http://github.com/alecjacobson/geometry-  
processing-mesh-reconstruction.git
```

Installation, Layout, and Compilation

See [introduction](#).

Execution


Once built, you can execute the assignment from inside the `build/` using

```
./mesh-reconstruction [path to point cloud]
```


Background

In this assignment, we will be implementing a simplified version of the method in "Poisson Surface Reconstruction" by Kazhdan et al. 2006. (Your first "task" will be to read and understand this paper).

Many scanning technologies output a set of n point samples \mathbf{P} on the surface of the object in question. From these points and perhaps the location of the camera, one can also estimate normals \mathbf{N} to the surface for each point $\mathbf{p} \in \mathbf{P}$. This image shows the `data/elephant.pwn` input data with a white dot for each point and a red line segment pointing outward for each corresponding normal vector.



For shape analysis, visualization and other downstream geometry processing phases, we would like to convert this finitely sampled *point cloud* data into an *explicit continuous surface representation*: i.e., a **triangle mesh** (a special case of a **polygon mesh**). This image shows the corresponding output mesh for `data/elephant.pwn` input data above:




Voxel-based Implicit Surface

Converting the point cloud directly to a triangle mesh makes it very difficult to ensure that the mesh meets certain *topological* postconditions: i.e., that it is [manifold](#), [closed](#), and has a small number of [holes](#).

Instead we will first convert the point cloud *sampling representation* into a an *implicit surface representation*: where the unknown surface is defined as the *level-set* of some function $g : \mathbf{R}^3 \Rightarrow \mathbf{R}$ mapping all points in space to a scalar value. For example, we may define the surface ∂S of some solid, volumetric shape S to be all points $\mathbf{x} \in \mathbf{R}^3$ such that $g(\mathbf{x}) = \sigma$, where we may arbitrarily set $\sigma = \frac{1}{2}$.

$$\partial S = \{\mathbf{x} \in \mathbf{R}^3 | g(\mathbf{x}) = \sigma\}.$$

On the computer, it is straightforward *discretize* an implicit function. We define a regular 3D grid of *voxels* containing at least the *bounding box* of S . At each node in the grid $\mathbf{x}_{i,j,k}$ we store the value of the implicit function $g(\mathbf{x}_{i,j,k})$. This defines g *everywhere* in the grid via *trilinear interpolation*.



For example, consider a point $\mathbf{x} = (\frac{1}{3}, \frac{1}{4}, \frac{1}{2})$ lying in the middle of the bottom-most, front-most, left-most cell. We know the values at the eight corners. Trilinear interpolation can be understood as [linear interpolation](#) in the x -direction by $\frac{1}{3}$ on each x -axis-aligned edge, resulting in four values *living* on the same plane. These can then be linearly interpolated in the y direction by $\frac{1}{4}$ resulting in two points on the same line, and finally in the z direction by $\frac{1}{2}$ to get to our evaluation point $(\frac{1}{3}, \frac{1}{4}, \frac{1}{2})$.

An implicit surface stored as the level-set of a trilinearly interpolated grid can be *contoured* into a triangle mesh via the [Marching Cubes Algorithm](#). For the purposes of this assignment, we will treat this as a [black box](#). Instead, we focus on determining what values for g to store on the grid.

Characteristic functions of solids

We assume that our set of points \mathbf{P} lie on the surface $\partial\mathbf{S}$ of some physical [solid](#) object \mathbf{S} . This solid object must have some non-trivial volume that we can calculate *abstractly* as the integral of unit density over the solid:

$$\int_{\mathbf{S}} 1 \, dA.$$

We can rewrite this definite integral as an indefinite integral over all of \mathbf{R}^3 :

$$\int_{\mathbf{R}^3} \chi_{\mathbf{S}}(\mathbf{x}) \, dA,$$

by introducing the [characteristic function](#) of \mathbf{S} , that is *one* for points inside of the shape and *zero* for points outside of \mathbf{S} :

$$\chi_{\mathbf{S}}(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x} \in \mathbf{S} \\ 0 & \text{otherwise.} \end{cases}$$

Compared to typical [implicit surface functions](#), this function represents the surface $\partial\mathbf{S}$ of the shape \mathbf{S} as the *discontinuity* between the one values and the zero values. Awkwardly, the gradient of the characteristic function $\nabla\chi_{\mathbf{S}}$ is *not defined* along $\partial\mathbf{S}$.

One of the key observations made in [Kazhdan et al. 2006] is that the gradient of a infinitesimally [mollified](#) (smoothed) characteristic function:

1. points in the direction of the normal near the surface $\partial\mathbf{S}$, and
2. is zero everywhere else.

Our goal will be to use our points \mathbf{P} and normals \mathbf{N} to *optimize* an implicit function g over a regular grid, so that its gradient ∇g meets these two criteria. In that way, our g will be an approximation of the mollified characteristic function.

Poisson surface reconstruction

Or: how I learned to stop worrying and minimize squared Gradients

Let us start by making two assumptions:

1. we know how to compute ∇g at each node location $\mathbf{x}_{i,j,k}$, and
2. our input points \mathbf{P} all lie perfectly at grid nodes: $\exists \mathbf{x}_{i,j,k} = \mathbf{p}_\ell$.

We will find out these assumptions are not realistic and we will have to relax them (i.e., we **will not** make these assumptions in the completion of the tasks). However, it will make the following algorithmic description easier on the first pass.

If our points \mathbf{P} lie at grid points, then our corresponding normals \mathbf{N} also *live* at grid points. This leads to a very simple set of linear equations to define a function g with a gradient equal to the surface normal at the surface and zero gradient away from the surface:

$$\nabla g(\mathbf{x}_{i,j,k}) = \mathbf{v}_{i,j,k} := \begin{cases} \mathbf{n}_\ell & \text{if } \exists \mathbf{p}_\ell = \mathbf{x}_{i,j,k}, \\ \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} & \text{otherwise.} \end{cases}$$

This is a *vector-valued* equation. The gradients, normals and zero-vectors are three-dimensional (e.g., $\nabla g \in \mathbb{R}^3$). In effect, this is *three equations* for every grid node.

Since we only need a single number at each grid node (the value of g), we have *too many* equations.

Like many geometry processing algorithms confronted with such an **over determined**, we will *optimize* for the solution that best *minimizes* the error of equation:

$$\|\nabla g(\mathbf{x}_{i,j,k}) - \mathbf{v}_{i,j,k}\|^2.$$

We will treat the error of each grid location equally by minimizing the sum over all grid locations:

$$\min_{\mathbf{g}} \sum_i \sum_j \sum_k \frac{1}{2} \|\nabla g(\mathbf{x}_{i,j,k}) - \mathbf{v}_{i,j,k}\|^2,$$

where \mathbf{g} (written in boldface) is a vector of *unknown* grid-nodes values, where $g_{i,j,k} = g(\mathbf{x}_{i,j,k})$.

Part of the convenience of working on a regular grid is that we can use the **finite difference method** to approximate the gradient ∇g on the grid.

After revisiting [our assumptions](#), we will be able to compute approximations of the x -, y - and z -components of ∇g via a [sparse matrix](#) multiplication of a "gradient matrix" G and our vector of unknown grid values g . We will be able to write the minimization problem above in matrix form:

$$\min_g \frac{1}{2} \|Gg - v\|^2,$$

or equivalently after expanding the norm:

$$\min_g \frac{1}{2} g^T G^T G g - g^T G^T v + \text{constant},$$

This is a quadratic "energy" function of the variables of g , its minimum occurs when an infinitesimal change in g produces no change in the energy:

$$\frac{\partial}{\partial g} \frac{1}{2} g^T G^T G g - g^T G^T v = 0.$$

Applying this derivative gives us a *sparse* system of linear equations

$$G^T G g = G^T v.$$

We will assume that we can solve this using a black box sparse solver.

Now, let's revisit [our assumptions](#).

Gradients on a regular grid

The gradient of a function g in 3D is nothing more than a vector containing partial derivatives in each coordinate direction:

$$\nabla g(\mathbf{x}) = \begin{pmatrix} \frac{\partial g(\mathbf{x})}{\partial x} \\ \frac{\partial g(\mathbf{x})}{\partial y} \\ \frac{\partial g(\mathbf{x})}{\partial z} \end{pmatrix}.$$


We will approximate each partial derivative individually. Let's consider the partial derivative in the x direction, $\partial g(\mathbf{x})/\partial x$, and we will assume without loss of generality that what we derive applies *symmetrically* for y and z .

The partial derivative in the x -direction is a one-dimensional derivative. This couldn't be easier to do with finite differences. We approximate the derivative of the function g with respect to the x direction is the difference between the function evaluated at one grid node and at the grid node *before* it in the x -direction then divided by the spatial distance between adjacent nodes h (i.e., the grid step size):


$$\frac{\partial g(\mathbf{x}_{i-\frac{1}{2},j,k})}{\partial x} = \frac{g_{i,j,k} - g_{i-1,j,k}}{h},$$

where we use the index $i - \frac{1}{2}$ to indicate that this derivative in the x -direction lives on a **staggered grid** *in between* the grid nodes where the function values for g .


The following pictures show a 2D example, where g lives on the nodes of a 5×5 blue grid:



The partial derivatives of g with respect to the x -direction $\partial g(\mathbf{x})/\partial x$ live on a 4×5 green, staggered grid:



The partial derivatives of g with respect to the y -direction $\partial g(\mathbf{x})/\partial y$ live on a 5×4 yellow, staggered grid:



Letting $\mathbf{g} \in \mathbf{R}^{n_x n_y n_z \times 1}$ be column vector of function values on the *primary grid* (blue in the example pictures), we can construct a sparse matrix $\mathbf{D}^x \in \mathbf{R}^{(n_x-1)n_y n_z \times n_x n_y n_z}$ so that each row $\mathbf{D}_{i-\frac{1}{2},j,k}^x \in \mathbf{R}^{1 \times n_x n_y n_z}$ computes the partial derivative at the corresponding staggered grid location $\mathbf{x}_{i-\frac{1}{2},j,k}$. The ℓ th entry in that row receives a value only for neighboring primary grid nodes:

$$\mathbf{D}_{i-\frac{1}{2},j,k}^x(\ell) = \begin{cases} -1 & \text{if } \ell = i - 1 \\ 1 & \text{if } \ell = i \\ 0 & \text{otherwise} \end{cases}.$$

Indexing 3D arrays

Now, obviously in our code we cannot *index* the column vector \mathbf{g} by a triplet of numbers $\{i, j, k\}$ or the rows of \mathbf{D}^x by the triplet $i - \frac{1}{2}, j, k$. We will assume that $\mathbf{g}_{i,j,k}$ refers to $\mathbf{g}(i+j*n_x+k*n_y*n_x)$. Similarly, for the staggered grid subscripts $i - \frac{1}{2}, j, k$ we will assume that $\mathbf{D}_{i-\frac{1}{2},j,k}^x(\ell)$ refers to the matrix entry $\mathbf{D}_x(i+j*n_x+k*n_y*n_x, \ell)$, where the $i - \frac{1}{2}$ has been *rounded down*.

We can similarly build matrices \mathbf{D}^y and \mathbf{D}^z and *stack* these matrices vertically to create a gradient matrix \mathbf{G} :

$$\mathbf{G} = \begin{pmatrix} \mathbf{D}^x \\ \mathbf{D}^y \\ \mathbf{D}^z \end{pmatrix} \in \mathbf{R}^{((n_x-1)n_y n_z + n_x(n_y-1)n_z + n_x n_y(n_z-1)) \times n_x n_y n_z}$$

This implies that our vector \mathbf{v} of zeros and normals in our minimization problem should not *live* on the primary, but rather it, too, should be broken into x -, y - and z -components that live of their respective staggered grids:

$$\mathbf{v} = \begin{pmatrix} \mathbf{v}^x \\ \mathbf{v}^y \\ \mathbf{v}^z \end{pmatrix} \in \mathbf{R}^{((n_x-1)n_y n_z + n_x(n_y-1)n_z + n_x n_y(n_z-1)) \times 1}.$$

This leads to addressing our second assumption.

B-b-b-b-but the input normals might not be at grid node locations?

At this point, we would *actually* liked to have had that our input normals were given component-wise on the staggered grid. Then we could immediately stick them into \mathbf{v} . But this doesn't make much sense as each normal \mathbf{n}_ℓ *lives* at its associated point \mathbf{p}_ℓ , regardless of any grids.

To remedy this, we will distribute each component of each input normal \mathbf{n}_ℓ to \mathbf{v} at the corresponding staggered grid node location.

For example, consider the normal \mathbf{n} at some point $\mathbf{x}_{1, \frac{1}{4}, \frac{1}{2}}$. Conceptually, we'll think of the x -component of the normal n_x as floating in the staggered grid corresponding to \mathbf{D}^x , in between the eight staggered grid locations:

$$\mathbf{x}_{\frac{1}{2}, 0, 0}, \mathbf{x}_{1\frac{1}{2}, 0, 0}, \mathbf{x}_{\frac{1}{2}, 1, 0}, \mathbf{x}_{1\frac{1}{2}, 1, 0}, \mathbf{x}_{\frac{1}{2}, 0, 1}, \mathbf{x}_{1\frac{1}{2}, 0, 1}, \mathbf{x}_{\frac{1}{2}, 1, 1}, \text{ and } \mathbf{x}_{1\frac{1}{2}, 1, 1}$$

Each of these staggered grid nodes has a corresponding x value in the vector \mathbf{v}^x .

We will distribute n_x to these entries in \mathbf{v}^x by *adding* a partial amount of n_x to each. I.e.,

$$v_{\frac{1}{2}, 0, 0}^x = w_{\frac{1}{2}, 0, 0} \left(\mathbf{x}_{1, \frac{1}{4}, \frac{1}{2}} \right) n_x, v_{1\frac{1}{2}, 0, 0}^x = w_{1\frac{1}{2}, 0, 0} \left(\mathbf{x}_{1, \frac{1}{4}, \frac{1}{2}} \right) n_x, \dots, v_{1\frac{1}{2}, 1, 1}^x = w_{1\frac{1}{2}, 1, 1} \left(\mathbf{x}_{1, \frac{1}{4}, \frac{1}{2}} \right) n_x.$$

where $w_{i+\frac{1}{2},j,k}(\mathbf{p})$ is the trilinear interpolation *weight* associate with staggered grid node $\mathbf{x}_{i+\frac{1}{2},j,k}$ to interpolate a value at the point \mathbf{p} . The trilinear interpolation weights so that:

$$n_x = w_{\frac{1}{2},0,0}(\mathbf{x}_{1,\frac{1}{4},\frac{1}{2}}) v_{\frac{1}{2},0,0}^x + w_{1\frac{1}{2},0,0}(\mathbf{x}_{1,\frac{1}{4},\frac{1}{2}}) v_{1\frac{1}{2},0,0}^x + \dots + w_{1\frac{1}{2},1,1}(\mathbf{x}_{1,\frac{1}{4},\frac{1}{2}}) v_{1\frac{1}{2},1,1}^x.$$

Since we need to do these for the x -component of each input normal, we will assemble a sparse matrix $\mathbf{W}^x \in n \times (n_x - 1)n_y n_z$ that *interpolates* \mathbf{v}^x at each point \mathbf{p} :

$$(\mathbf{W}^x \mathbf{v}^x) \in \mathbf{R}^{n \times 1}$$

the transpose of \mathbf{W}^x is not quite its *inverse*, but instead can be interpreted as *distributing* values onto staggered grid locations where \mathbf{v}^x lives:

$$\mathbf{v}^x = (\mathbf{W}^x)^\top \mathbf{N}^x.$$

Using this definition of \mathbf{v}^x and analogously for \mathbf{v}^y and \mathbf{v}^z we can construct the vector \mathbf{v} in our energy minimization problem above.

BTW, what's **Poisson** got to do with it?

The discrete energy minimization problem we've written looks like the squared norm of some gradients. An analogous energy in the smooth world is the **Dirichlet energy**:

$$E(g) = \int_{\Omega} \|\nabla g\|^2 dA$$

to *minimize* this energy with respect to g as an unknown *function*, we need to invoke **Calculus of Variations** and **Green's First Identity**. In doing so we find that minimizers will satisfy:

$$\nabla \cdot \nabla g = 0 \text{ on } \Omega,$$

known as **Laplace's Equation**.

If we instead start with a slightly different energy:

$$E(g) = \int_{\Omega} \|\nabla g - V\|^2 dA,$$

where V is a vector-valued function. Then applying the same machinery we find that minimizers will satisfy:

$$\nabla \cdot \nabla g = \nabla \cdot V \text{ on } \Omega,$$

known as **Poisson's equation**.

Notice that if we interpret the transpose of our gradient matrix \mathbf{G}^T as a *divergence matrix* (we can and we should), then the structure of these smooth energies and equations are directly preserved in our discrete energies and equations.

This kind of *structure preservation* is a major criterion for judging discrete methods.

Choosing a good iso-value

Constant functions have no gradient. This means that we can add a constant function to our implicit function g without changing its gradient:

$$\nabla g = \nabla(g + c) = \nabla g + \nabla c = \nabla g + 0.$$

The same is true for our discrete gradient matrix \mathbf{G} : if the vector of grid values \mathbf{g} is constant then $\mathbf{G}\mathbf{g}$ will be a vector zeros.

This is potentially problematic for our least squares solve: there are many solutions, since we can just add a constant. Fortunately, we *don't really care*. It's elegant to say that our surface is defined at $g = 0$, but we'd be just as happy declaring that our surface is defined at $g = c$.

To this end we just need to find a *solution* \mathbf{g} , and then to pick a good iso-value σ .

As suggested in [Kazhdan et al. 2006], we can pick a good iso-value by interpolating our solution \mathbf{g} at each of the input points (since we know they're on the surface) and averaging their values. For an appropriate interpolation matrix \mathbf{W} on the *primary (non-staggered) grid* this can be written as:

$$\sigma = \frac{1}{n} \mathbf{1}^T \mathbf{W} \mathbf{g},$$

where $\mathbf{1} \in \mathbf{R}^{n \times 1}$ is a vector of ones.

Just how much does this assignment simplify [Kazhdan et al. 2006]?

Besides the insights above, a major contribution of [Kazhdan et al. 2006] was to setup and solve this problem on an [adaptive grid](#) rather than a regular grid. They also track "confidence" of their input data effecting how they smooth and interpolate values. As a result, their method is one of the most highly used surface reconstruction techniques to this day.

Consider adding your own insights to the wikipedia entry for [this method](#).

Tasks

Read [Kazhdan et al. 2006]

This reading task is not directly graded, but it's expected that you read and understand this paper before moving on to the other tasks.

src/fd_interpolate.cpp

Given a regular finite-difference grid described by the number of nodes on each side (n_x , n_y and n_z), the grid spacing (h), and the location of the bottom-left-front-most corner node ($corner$), and a list of points (P), construct a sparse matrix W of trilinear interpolation weights so that $P = W * X$.

src/fd_partial_derivative.cpp

Given a regular finite-difference grid described by the number of nodes on each side (n_x , n_y and n_z), the grid spacing (h), and a desired direction, construct a sparse matrix D to compute first partial derivatives in the given direction onto the *staggered grid* in that direction.

src/fd_grad.cpp

Given a regular finite-difference grid described by the number of nodes on each side (n_x , n_y and n_z), and the grid spacing (h), construct a sparse matrix G to compute gradients with each component on its respective staggered grid.

Hint

Use `fd_partial_derivative` to compute D_x , D_y , and D_z and then simply concatenate these together to make G .

src/poisson_surface_reconstruction.cpp

Given a list of points P and the list of corresponding normals N , construct an implicit function g over a regular grid (*built for you*) using approach described above.

You will need to *distribute* the given normals N onto the staggered grid values in v via sparse trilinear interpolation matrices w_x , w_y and w_z for each staggered grid.

Then you will need to construct and solve the linear system $G^T G g = G^T v$.

Determine the iso-level σ to extract from the g .

Feed this implicit function `g` to `igl::copyleft::marching_cubes` to contour this function into a triangle mesh `v` and `F`.

Make use of `fd_interpolate` and `fd_grad`.

Hint

Eigen has many different [sparse matrix solvers](#). For these *very regular* matrices, it seems that the [conjugate gradient method](#) will outperform direct methods such as [Cholesky factorization](#). Try `Eigen::BiCGSTAB`.

Hint

Debug in debug mode with assertions enabled. For Unix users on the command line use:

```
cmake -DCMAKE_BUILD_TYPE=Debug ../
```

but then try out your code in *release mode* for much better performance

```
cmake -DCMAKE_BUILD_TYPE=Release ../
```

Releases

No releases published

Packages

No packages published

Contributors 3



alecjacobson Alec Jacobson



psarahdactyl Sarah Kushner

Languages

Pawn 99.8%

Other 0.2%